

MySQL主从架构

一、背景

随着公司业务的发展，对数据库的要求越来越高，以前的单机MySQL肯定是玩不转了，亟需升级成MySQL集群，这是一个公司在业务发展时不得不面临的问题。就单机MySQL而言，自己玩玩可以，运用到实际项目中，那肯定要挨批的。一方面数据不安全，万一数据库的电脑磁盘坏了，就坑了。另一方面数据库的并发能力是有限的，一般并发数200~500就差不多了，当然你要继续往上增，也是可以的，那就影响整体Mysql的响应时间。

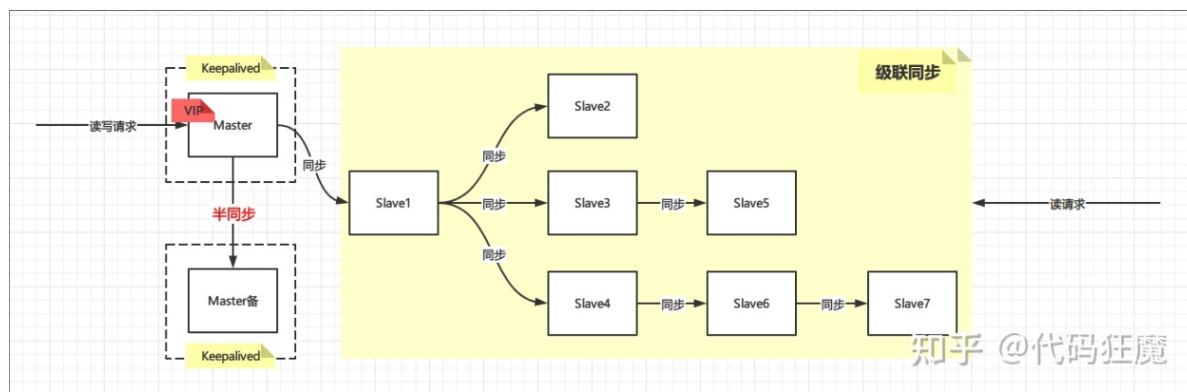
二、集群方案

2.1 主从数据同步（主写从读）

具体配置参照本文档下的《主从架构配置》文件

该配置只能用于主写从读，如果向从数据库写入数据，由于不能将数据同步到主数据库，从而导致主从数据库不一致的问题。

2.2 普通双机热备



优点：

- 读写分离，增加整体性能
- Master有故障转移（采用Keepalived实现），当Master宕机时，备用Master自动顶上去，且IP无变化（VIP）
- Master和备用Master之间采用半同步机制，最大程度上保证数据一致性

缺点：

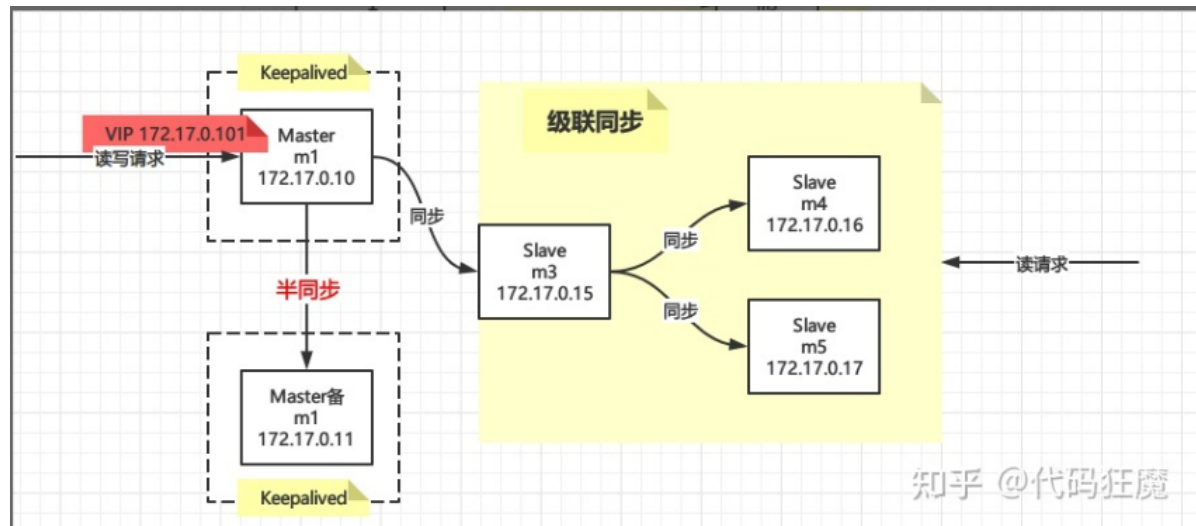
- Master故障转移时，Slave需要用脚本手动切换，增加复杂度
- 数据存在一致性问题，因为异步，所以Slave的数据一定不是最新的，需要等待一个时间窗后才能读取
- Slave过多时Slave对Master的负载以及网络带宽都会成为一个严重的问题

2.3 Keepalived双主热备

2.3.1 架构

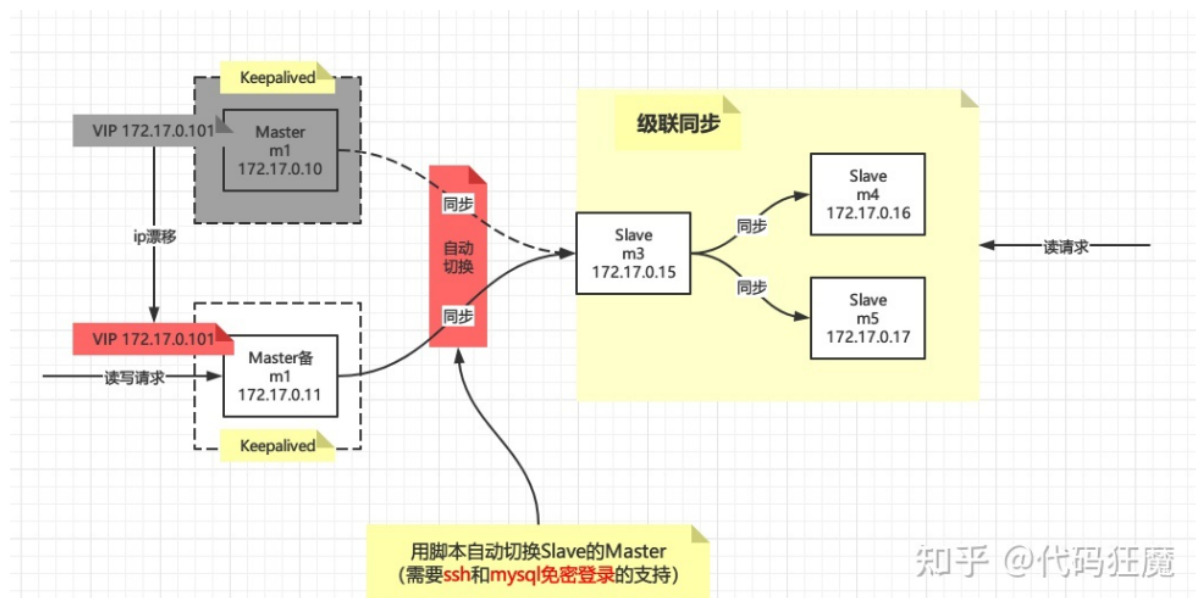
m1和m2作为双主，**互为主从**且同步方式为**半同步**，为什么呢？因为这样在Master主机宕机恢复，对Master备机写入的数据也能重新同步到Master主，此时的主备角色正好颠倒过来，所以要设置互为主从且半同步。

m3作为级联复制的中间节点，新增加m4和m5作为新的Slave，正常情况下架构如下：



当发生故障时，VIP自动漂移到Master备机上面，且通过Keepalived的状态脚本自动切换Slave，即m3到Master备机上，当Master主机（m1）恢复时，VIP又会自动漂移到m1上并且m3也会自动切换到m1上！

Keepalived状态切换脚本需要ssh和mysql免密登录的支持，其原理就是m1或者m2检测到自己成为Master角色时，通过ssh登录到m3上执行MySQL的切换命令



此文档只实现了Keepalived主备切换功能（即m1与m2的切换），如需切实现slave由m1切换到m2，请参照其他文档，添加相应脚本。

2.3.2 环境准备

Mysql 版本: 8.0.27

```
docker run -d --name m1 -p 3306:3306 -e MYSQL_ROOT_PASSWORD=123456 -e TZ="Asia/Shanghai" -v /home/rebei/m1/config:/etc/mysql/conf.d mysql --authentication_policy=mysql_native_password

docker restart m1
docker exec -it m1 mysql -uroot -p123456
-----

docker run -d --name m2 -p 3306:3306 -e MYSQL_ROOT_PASSWORD=123456 -e TZ="Asia/Shanghai" -v /home/rebei/m2/config:/etc/mysql/conf.d mysql --authentication_policy=mysql_native_password

docker restart m2
docker exec -it m2 mysql -uroot -p123456
-----

#-----
#docker run -d --name s1 -p 4306:3306 -e MYSQL_ROOT_PASSWORD=123456 -e TZ="Asia/Shanghai" -v #/home/rebei/s1/config:/etc/mysql/conf.d mysql --authentication_policy=mysql_native_password

#docker restart s1
#docker exec -it s1 mysql -uroot -p123456
-----
```

注意: 需要确认一下每台设备上的server_uuid是否相同, 如果相同, 则需重置, 否则后面的主从复制会出错。

```
docker exec -it m1 /bin/sh
cat /var/lib/mysql/auto.cnf
#如果一致, 则删除重新启动服务
rm -f /var/lib/mysql/auto.cnf
docker restart m1
```

2.3.3 双主配置

Master主机: 192.168.200.104

Master备机: 192.168.200.105

- Master主机 (m1)

```
vim /home/rebei/m1/config/my.cnf
# 修改配置
[mysqld]
user=mysql
character-set-server=utf8mb4
collation-server=utf8mb4_general_ci
```

```
log-bin=mysql-bin-master
server_id=1
log_replica_updates=1
plugin-
load="rpl_semi_sync_source=semisync_source.so;rpl_semi_sync_replica=semisync_replica.so"
rpl_semi_sync_source_enabled=1
rpl_semi_sync_replica_enabled=1

[client]
default-character-set=utf8mb4
```

- Master备机 (m2)

说明: server_id必须与m1不同

```
vim /home/rebei/m2/config/my.cnf
# 修改配置
[mysqld]
user=mysql
character-set-server=utf8mb4
collation-server=utf8mb4_general_ci
log-bin=mysql-bin-master
server_id=2
log_replica_updates=1
plugin-
load="rpl_semi_sync_source=semisync_source.so;rpl_semi_sync_replica=semisync_replica.so"
rpl_semi_sync_source_enabled=1
rpl_semi_sync_replica_enabled=1

[client]
default-character-set=utf8mb4
```

可使用show plugins查看主从插件是否安装成功

```
show plugins
```

show plugins 输入一个 SQL 表达式来过滤结果 (使用 Ctrl+Space)					
	ABC Name	ABC Status	ABC Type	ABC Library	ABC License
36	MyISAM	ACTIVE	STORAGE ENGINE	[NULL]	GPL
37	MRG_MYISAM	ACTIVE	STORAGE ENGINE	[NULL]	GPL
38	PERFORMANCE_SCHEMA	ACTIVE	STORAGE ENGINE	[NULL]	GPL
39	TempTable	ACTIVE	STORAGE ENGINE	[NULL]	GPL
40	ARCHIVE	ACTIVE	STORAGE ENGINE	[NULL]	GPL
41	BLACKHOLE	ACTIVE	STORAGE ENGINE	[NULL]	GPL
42	FEDERATED	DISABLED	STORAGE ENGINE	[NULL]	GPL
43	ngram	ACTIVE	FTPARSER	[NULL]	GPL
44	mysqlx_cache_cleaner	ACTIVE	AUDIT	[NULL]	GPL
45	mysqlx	ACTIVE	DAEMON	[NULL]	GPL
46	rpl_semi_sync_source	ACTIVE	REPLICATION	semisync_source	GPL
47	rpl_semi_sync_replica	ACTIVE	REPLICATION	semisync_replica	GPL

- 创建同步用户

在m1与m2机器的MySQL上执行一下脚本

```
DROP USER 'rep1'@'%';
CREATE USER 'rep1'@'%' IDENTIFIED BY 'rep1';
GRANT REPLICATION SLAVE ON *.* TO 'rep1'@'%';
flush privileges;
```

查看Master状态

```
SHOW MASTER STATUS;
```

SHOW MASTER STATUS 输入一个 SQL 表达式来过滤结果 (使用 Ctrl+Space)					
	ABC File	123 Position	ABC Binlog_Do_DB	ABC Binlog_Ignore_DB	ABC Executed_Gtid_Set
1	mysql-bin-master.000001	156			
	binglog文件	偏移量			

2.3.4 互为主从

- (1) 将Master主挂载成为Master备机的从机

```
STOP SLAVE;
RESET SLAVE;

CHANGE MASTER TO
MASTER_HOST='192.168.200.105',
MASTER_USER='rep1',
MASTER_PASSWORD='rep1',
MASTER_LOG_FILE='mysql-bin-master.000001',
MASTER_LOG_POS=156;

START SLAVE;
```

说明:

MASTER_HOST: 主机IP (该sql脚本是在从机上执行)

MASTER_LOG_FILE: 主机对应的偏移文件

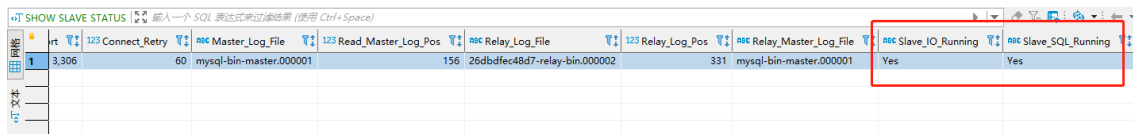
MASTER_LOG_POS: 主机对应的偏移量

(2) 将Master备机挂载成为Master主机的从机，需修改对应的IP及偏移量

```
STOP SLAVE;  
RESET SLAVE;  
  
CHANGE MASTER TO  
MASTER_HOST='192.168.200.104',  
MASTER_USER='repl',  
MASTER_PASSWORD='repl',  
MASTER_LOG_FILE='mysql-bin-master.000001',  
MASTER_LOG_POS=156;  
  
START SLAVE;
```

- 查看下是否挂载成功

```
SHOW SLAVE STATUS;
```



Slave_IO_Running	Slave_SQL_Running
Yes	Yes

- 查看半同步是否开启

```
show variables like 'rpl_semi_sync_source_enabled';  
show variables like 'rpl_semi_sync_replica_enabled';
```

最后在自行测试下载Master主机上面写数据Master备机能否同步，反过来再测试下。

2.3.5 配置keepalived

(1) 安装keepalived

```
yum install -y keepalived
```

注意：Keepalived就安装在Docker运行的宿主机上面，安装在docker里面IP无法绑定(亲测)

(2) Master主机配置

```
vi /etc/keepalived/keepalived.conf
```

```
! Configuration File for keepalived  
  
global_defs {  
    router_id HA-M1  
    script_user root  
    enable_script_security  
}
```

```

vrrp_script chk_mysql {
    script /etc/keepalived/chk_mysql.sh
    interval 60
    fall 3
    rise 2
}

vrrp_instance VI_MYSQL {
    state MASTER
    interface ens33
    virtual_router_id 100
    priority 100
    advert_int 1

    authentication {
        auth_type PASS
        auth_pass 1111
    }

    virtual_ipaddress {
        192.168.200.200 dev ens33 label ens33:vip
    }

    track_script {
        chk_mysql
    }
}

```

- router_id: 标识, 两台需不一样
- state MASTER: 表明这台是Master角色
- priority: 优先级, MASTER角色比BACKUP高!
- virtual_router_id: 虚拟路由编号, 两台需要一致
- 192.168.200.200 dev ens33 label ens33:vip: 这就是我们配置的VIP: 192.168.200.200
- script /etc/keepalived/chk_mysql.sh: MySQL的检测脚本, 定时去检测MySQL进程是否挂掉, 如果挂掉, 在脚本里面重启之, 定时通过interval配置, 上面配置的是2s检测一次

注意: interval如果过小可能会导致如下错误: /etc/[keepalived](#)/chk_nginx.sh exited due to signal 15

(3) Master备机配置

```
vi /etc/keepalived/keepalived.conf
```

! Configuration File for keepalived

```

global_defs {
    router_id HA-M2  #router id配置,与主机不一致

```

```

script_user root
enable_script_security
}

vrrp_script chk_mysql { #定义MYSQL状态检测脚本
    script /etc/keepalived/chk_mysql.sh
    interval 30
    fall 3
    rise 2
}

vrrp_instance VI_MYSQL {
    state BACKUP      #角色配置
    interface ens33   #网卡配置
    virtual_router_id 100 #必须与主机一致
    nopreempt
    priority 50       #优先级配置，必须小于MASTER
    advert_int 1
    # mcast_src_ip 192.168.200.104

    authentication {
        auth_type PASS
        auth_pass 1111
    }

    virtual_ipaddress {
        192.168.200.200 dev ens33 label ens33:vip #虚拟IP
    }

    track_script {
        chk_mysql
    }
}

```

(3) 创建Mysql状态检测脚本

```

touch /etc/keepalived/chk_mysql.sh
chmod +x /etc/keepalived/chk_mysql.sh

```

chk_mysql.sh

```

cat > /etc/keepalived/chk_mysql.sh <<EOF
#!/bin/bash
MYSQL=/usr/bin/mysql
MYSQL_HOST=192.168.3.254
MYSQL_USER=root
MYSQL_PASSWORD=123456
CHECK_TIME=3

#mysql is working MYSQL_OK is 1 , mysql down MYSQL_OK is 0

```



```

MYSQL_OK=1

function check_mysql_health () {
    $MYSQL -h $MYSQL_HOST -u$MYSQL_USER -p${MYSQL_PASSWORD} -e "show
status;" >/dev/null 2>&1
    if [ $? = 0 ] ; then
        MYSQL_OK=1
    else
        MYSQL_OK=0
    fi
    return $MYSQL_OK
}
while [ $CHECK_TIME -ne 0 ]
do
    let "CHECK_TIME -= 1"
    check_mysql_health
    if [ $MYSQL_OK = 1 ] ; then
        CHECK_TIME=0
        exit 0
    fi
    if [ $MYSQL_OK -eq 0 ] && [ $CHECK_TIME -eq 0 ]
    then
        systemctl stop keepalived
        exit 1
    fi
    sleep 1
done
EOF

```

(4) 启动keepalived，并查看虚拟IP是否绑定成功

```

#实时查看keepalived启动日志，看是否有错误
tail -f /var/log/messages
#启动keepalived
systemctl start keepalived
#查看虚拟IP是否绑定成功
ifconfig

```

```
[root@localhost ~]# ifconfig
br-3869e3e8a3ac: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.18.0.1 netmask 255.255.0.0 broadcast 172.18.255.255
    ether 02:42:3b:19:40:84 txqueuelen 0 (Ethernet)
    RX packets 246 bytes 81514 (79.6 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 301 bytes 38148 (37.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

docker0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:f1ff:fe01:9973 prefixlen 64 scopeid 0x20<link>
    ether 02:42:f1:01:99:73 txqueuelen 0 (Ethernet)
    RX packets 246 bytes 81514 (79.6 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 301 bytes 38148 (37.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.200.105 netmask 255.255.255.0 broadcast 192.168.200.255
    inet6 fe80::b0f4:3b81:4994:3b64 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:84:36:2c txqueuelen 1000 (Ethernet)
    RX packets 4383 bytes 320223 (312.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 7046 bytes 870802 (850.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33:vip: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.200.200 netmask 255.255.255.255 broadcast 0.0.0.0
    ether 00:0c:29:84:36:2c txqueuelen 1000 (Ethernet)

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 388 bytes 25120 (24.5 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 388 bytes 25120 (24.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

veth54817f6: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::44be:30ff:fe18:1910 prefixlen 64 scopeid 0x20<link>
    ether 46:be:30:18:19:10 txqueuelen 0 (Ethernet)
    RX packets 246 bytes 84958 (82.9 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 309 bytes 38804 (37.8 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

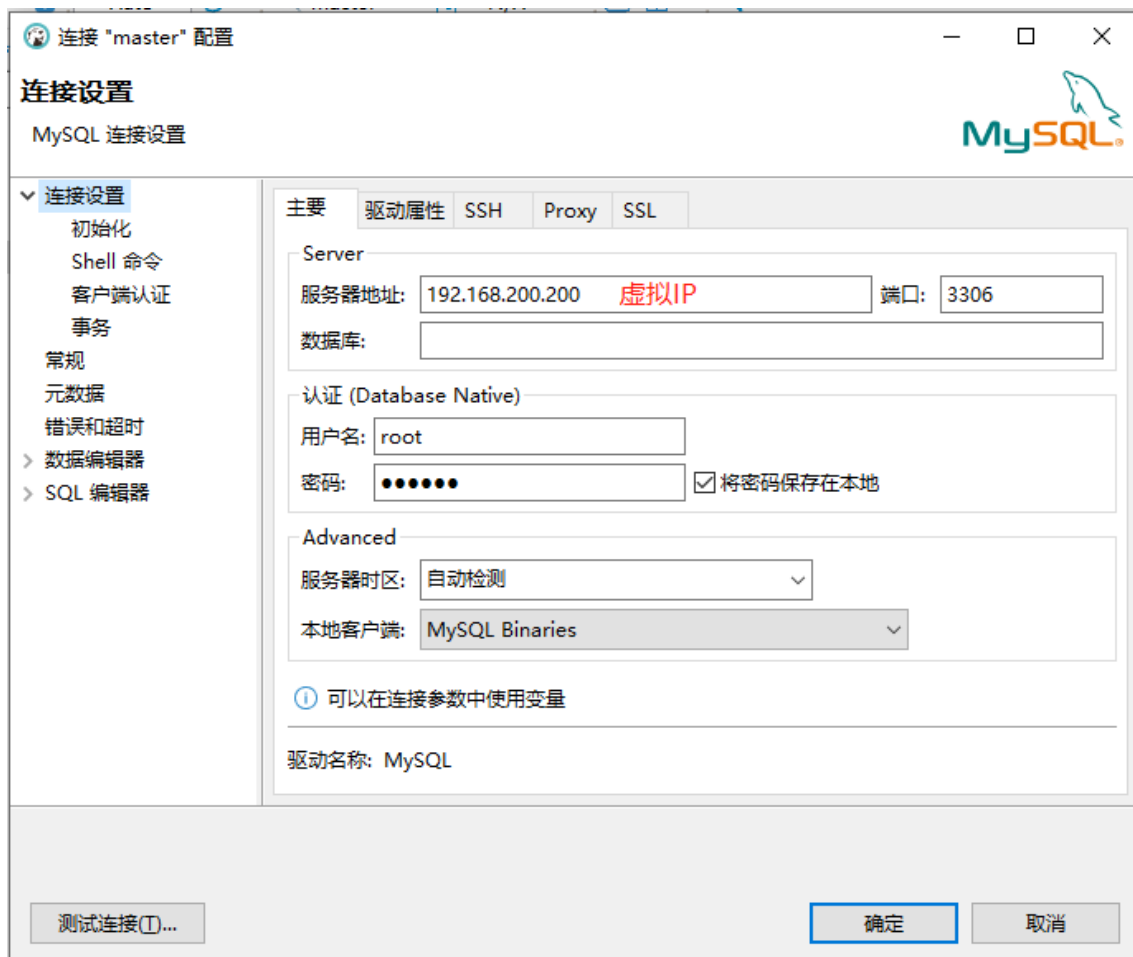
(5) 安装Mysql客户端 (也可直接通过Windows连接)

```
wget https://dev.mysql.com/get/mysql80-community-release-el7-5.noarch.rpm
rpm -ivh mysql80-community-release-el7-5.noarch.rpm
yum install mysql-community-client
```

2.3.6 测试

- 确认是否可以通过虚拟IP连接MYSQL

先分别确认m1与m2容器中的MYSQL是否可以正常访问, 再是否可以确认虚拟IP访问MYSQL



- 查看keepalived是否实现主备切换功能

Master主机

```
#关闭Master主机的MySQL服务
docker stop m1
查看keepalived状态
systemctl status keepalived
#查看虚拟IP是否绑定情况（V此时IP已经漂移到备机）
ifconfig
```

```

[root@www ~]# docker stop m1
m1
[root@www ~]# systemctl status keepalived
● keepalived.service - LVS and VRRP High Availability Monitor
   Loaded: loaded (/usr/lib/systemd/system/keepalived.service; disabled; vendor preset: disabled)
   Active: inactive (dead)

Jun 27 00:43:10 www.yptest.com Keepalived_vrrp[2960]: Sending gratuitous ARP on ens33 for 192.168.200.200
Jun 27 00:43:10 www.yptest.com Keepalived_vrrp[2960]: Sending gratuitous ARP on ens33 for 192.168.200.200
Jun 27 00:43:59 www.yptest.com Keepalived[2958]: Stopping
Jun 27 00:43:59 www.yptest.com systemd[1]: Stopping LVS and VRRP High Availability Monitor...
Jun 27 00:43:59 www.yptest.com Keepalived_vrrp[2960]: VRRP_Instance(VI_MYSQL) sent 0 priority
Jun 27 00:43:59 www.yptest.com Keepalived_vrrp[2960]: VRRP_Instance(VI_MYSQL) removing protocol VIPs.
Jun 27 00:43:59 www.yptest.com Keepalived_healthcheckers[2959]: Stopped
Jun 27 00:44:00 www.yptest.com Keepalived_vrrp[2960]: Stopped
Jun 27 00:44:00 www.yptest.com Keepalived[2958]: Stopped Keepalived v1.3.5 (03/19,2017), git commit v1.3.5-6-g6fa32f2
Jun 27 00:44:00 www.yptest.com systemd[1]: Stopped LVS and VRRP High Availability Monitor.
[root@www ~]# ifconfig
br-033be864c800: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.19.0.1 netmask 255.255.0.0 broadcast 172.19.255.255
    ether 02:42:30:3c:5e:76 txqueuelen 0 (Ethernet)
    RX packets 3434 bytes 267656 (261.3 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4151 bytes 639262 (624.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-a2234a6e04d5: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.20.0.1 netmask 255.255.0.0 broadcast 172.20.255.255
    ether 02:42:64:7d:f3:7b txqueuelen 0 (Ethernet)
    RX packets 3434 bytes 267656 (261.3 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4151 bytes 639262 (624.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

br-be1e61a4ceb5: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.18.0.1 netmask 255.255.0.0 broadcast 172.18.255.255
    ether 02:42:c3:d4:4c:bc txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:6dff:fe58:54c prefixlen 64 scopeid 0x20<link>
    ether 02:42:6d:58:05:4c txqueuelen 0 (Ethernet)
    RX packets 438 bytes 168544 (164.5 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 530 bytes 64070 (62.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.200.104 netmask 255.255.255.0 broadcast 192.168.200.255
    inet6 fe80::c256:8d20:5d38:7ddc prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:ed:4c:c1 txqueuelen 1000 (Ethernet)
    RX packets 3434 bytes 267656 (261.3 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4151 bytes 639262 (624.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

Master备机

ifconfig

```
[root@localhost ~]# ifconfig
br-3869e3e8a3ac: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.18.0.1 netmask 255.255.0.0 broadcast 172.18.255.255
    ether 02:42:3b:19:40:84 txqueuelen 0 (Ethernet)
    RX packets 341 bytes 89534 (87.4 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 390 bytes 48064 (46.9 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

docker0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:f1ff:fe01:9973 prefixlen 64 scopeid 0x20<link>
    ether 02:42:f1:01:99:73 txqueuelen 0 (Ethernet)
    RX packets 341 bytes 89534 (87.4 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 390 bytes 48064 (46.9 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.200.105 netmask 255.255.255.0 broadcast 192.168.200.255
    inet6 fe80::b0f4:3b81:4994:3b64 prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:84:36:2c txqueuelen 1000 (Ethernet)
    RX packets 6396 bytes 457614 (446.8 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 10630 bytes 1257398 (1.1 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens33:vip: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.200.200 netmask 255.255.255.255 broadcast 0.0.0.0
    ether 00:0c:29:84:36:2c txqueuelen 1000 (Ethernet)

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 524 bytes 33280 (32.5 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 524 bytes 33280 (32.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```